

# Improved Gradient-Based Neural Architecture Search

Zhexuan HUANG

Supervised by: Bruno CONCHE

July 10, 2019


# Overview

- 1 Introduction
- 2 State-of-the-art Neural Architecture Search (NAS) strategies
- 3 NAS with unbiased and low variance gradient estimators
- 4 Experiments
- 5 Conclusion

# Introduction

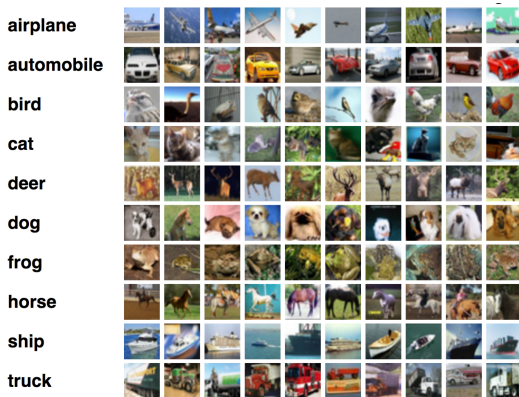
# Motivation of Neural Architecture Search (NAS)

- Designing neural network architectures is hard
- A lot of intuition and possibilities to design them
- Can we learn good architectures automatically?



# Image classification over CIFAR10

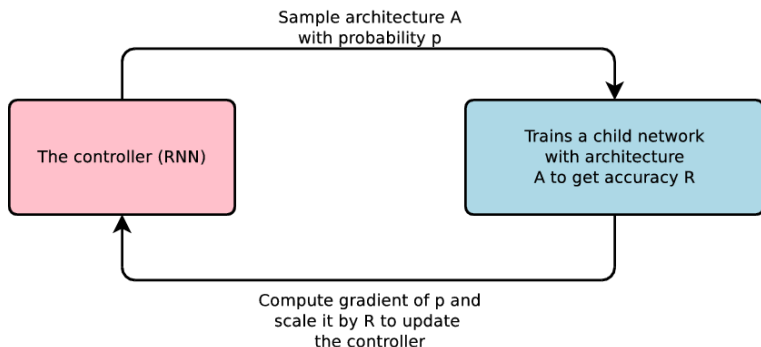
- CIFAR-10 dataset consists of 60,000 32x32 colour images in 10 classes
- To have a better comparison since many NAS researches use CIFAR10 as dataset



# State-of-the-art Neural Architecture Search (NAS) strategies

# Original Neural Architecture Search (NAS)

- Architecture Search algorithm proposed by Google in 2016 (Zoph and Le, 2016, Neural Architecture Search with Reinforcement Learning)
- Use un controller(RNN) to generate features of Neural Networks
- Train each NN to converge and return its accuracy to controller as rewards, update controller with reinforce policy



# Result of NAS

## Author's implementation

Used 1800 GPU days to search a model with 2.65% error rates over CIFAR10

## Analysis

- Reinforcement Learning(RL) based algorithm
- In each controller loop, need to train a child network to converge to get accuracy R
- Computational expensive and time consuming



# List of different state-of-the-art NAS strategies

## Reinforcement Learning based

- NAS(Zoph and Le, 2016), 1800 GPU days
- ENAS(Pham et al., 2018), 0.5 GPU days

## Evolution based

- AmoebaNet(Real et al., 2019), 3150 GPU days
- Hierarchical Evolution(H. Liu, Simonyan, Vinyals, et al., 2017), 300 GPU days

## Bayesian optimization

- PNAS(C. Liu et al., 2018), 225 GPU days

## Gradient based

- DARTS(H. Liu, Simonyan, and Yang, 2019), 1 GPU days
- SNAS(Xie et al., 2019), 1.5 GPU days

# Search space of Gradient based NAS

- Search space: set of all candidate architectures
- The same search space as in NAS, ENAS, DARTS and SNAS
- Search for computation cells, which is represented by a directed acyclic graph(DAG)
- Each edge  $(i, j)$  in DAG is associated with some operation  $\tilde{O}_{i,j}$ . The possible choice for  $\tilde{O}_{i,j}$  is given in priority (e.g. Conv  $3 \times 3$ , Maxpool  $3 \times 3$ , None, Identity, etc).




Figure: DAG, image retrieved from (H. Liu, Simonyan, and Yang, 2019)

# Search space of Gradient based NAS

- Two kinds of computation cells: normal cell and reduction cell
- Whole architecture is obtained by stacking two kinds of cell together






Figure: A conceptual of global structure.

# Examples of cells



(a) Normal cell found by SNAS



(b) Reduction cell found by SNAS

**Figure:** Normal cell and reduction cell (child graph) found by SNAS on CIFAR-10 (Xie et al., 2019) (a) Normal cell. (b) Reduction cell.

# Reformulate of problem with stochastic modeling

- As defined in SNAS, each  $\tilde{O}_{i,j}$  is probabilistic:

$$\tilde{O}_{i,j} = \begin{cases} O_{i,j}^1 & \text{with probability } p_{i,j}^1 = \frac{\exp(\alpha_{i,j}^1)}{\sum_{k=1}^A \exp(\alpha_{i,j}^k)} \\ \dots & \\ O_{i,j}^A & \text{with probability } p_{i,j}^A = \frac{\exp(\alpha_{i,j}^A)}{\sum_{k=1}^A \exp(\alpha_{i,j}^k)}, \end{cases}$$

where  $\alpha = (\alpha_{i,j}^k)_{i < j, k=1, \dots, A}$  are architecture parameters.

- Equivalent to  $\tilde{O}_{i,j} = Z_{i,j} O_{i,j}$  where

$$Z_{i,j} = \begin{cases} [1, 0, \dots, 0] & \text{with probability } p_{i,j}^1 = \frac{\exp(\alpha_{i,j}^1)}{\sum_{k=1}^A \exp(\alpha_{i,j}^k)} \\ \dots & \\ [0, 0, \dots, 1] & \text{with probability } p_{i,j}^A = \frac{\exp(\alpha_{i,j}^A)}{\sum_{k=1}^A \exp(\alpha_{i,j}^k)}, \end{cases}$$

# Objective

- Two kinds of parameters: architecture parameters  $\alpha$  and operation parameters  $\theta$
- Objective is to minimize the expected loss:

$$\min_{\alpha, \theta} \mathbb{E}_{Z \sim p_{\alpha}(Z)} [L_{\theta}(Z)], \quad (1)$$

which can be estimated by Monte-Carlo sampling.

- A conceptual of calculation loss retrieved from Xie et al., 2019

# Gradient estimators: Overview

- The objective  $\min_{\alpha, \theta} \mathbb{E}_{Z \sim p_{\alpha}(Z)}[L_{\theta}(Z)]$  can be optimized with gradient-descent algorithm
- Impossible to calculate the exact gradients
- $\nabla_{\theta} \mathbb{E}_{Z \sim p_{\alpha}(Z)}[L_{\theta}(Z)] = \mathbb{E}_{Z \sim p_{\alpha}(Z)}[\nabla_{\theta} L_{\theta}(Z)]$  can be estimated by Monte-Carlo sampling
- Different techniques to estimate  $\nabla_{\alpha} \mathbb{E}_{Z \sim p_{\alpha}(Z)}[L_{\theta}(Z)]$

# Gradient estimators: Score Function (SF) Estimators

Also known as REINFORCE(J.Williams, 1992), based on identity:

$$\nabla_{\alpha} \mathbb{E}_{Z \sim p_{\alpha}(Z)}[L_{\theta}(Z)] = \mathbb{E}_{Z \sim p_{\alpha}(Z)}[L_{\theta}(Z) \nabla_{\alpha} \log p_{\alpha}(Z)]. \quad (2)$$

## Analysis

- Unbiased estimator
- Depends only on the final result of  $L_{\theta}(Z)$
- Not require to calculate the back-propagation
- Extremely high variance



# Gradient estimators: Reparameterization Trick

If  $Z$  can be rewrite as  $Z = g(\alpha, \epsilon)$  where  $\epsilon \sim p_\epsilon$ , then

$$\nabla_\alpha \mathbb{E}_{Z \sim p_\alpha(Z)} [L_\theta(Z)] = \mathbb{E}_{\epsilon \sim p_\epsilon} [L'_\theta(g(\alpha, \epsilon)) \nabla_\alpha g(\alpha, \epsilon)]. \quad (3)$$

## Example

If  $X \sim \mathcal{N}(\mu, \sigma^2)$  then  $\nabla_\sigma \mathbb{E}[f(X)] = \mathbb{E}[f'(\mu + \sigma N)N]$  for  $N \sim \mathcal{N}(0, 1)$ .

## Analysis

- Unbiased, low variance, better than SF
- Only applicable for random variable which is reparameterizable
- Need to calculate the back-propagation

# Gradient estimators: SNAS and Gumbel-Softmax

Main difficulty to use reparameterization trick

$Z$  in our case is discrete random variable and is not reparameterizable.

Solution proposed in SNAS

- Relax  $Z$  to  $Z \rightarrow \tilde{Z}$  by Gumbel-Softmax, where

$$\tilde{Z}_{i,j}^k = g_{i,j}(\alpha, U) = \frac{\exp((\alpha_{i,j}^k - \log(-\log(U_{i,j}^k)))/\lambda)}{\sum_{l=1}^A \exp((\alpha_{i,j}^l - \log(-\log(U_{i,j}^l)))/\lambda)}, \quad (4)$$

$U = \{U_{i,j}^k\}_{i,j,k}$  are some independent uniform random variables,  $\lambda$  is the temperature of the Gumbel softmax, which is annealed to zero in SNAS.

- Minimize the approximated loss function with reparameterization:

$$\min_{\alpha, \theta} \mathbb{E}_{\tilde{Z} \sim \tilde{p}_\alpha(\tilde{Z})} [L_\theta(\tilde{Z})] = \min_{\alpha, \theta} \mathbb{E}_U [L_\theta(g(\alpha, U))], \quad (5)$$

# Gradient estimators: Intuition of Gumbel-Softmax

## Proposition 1

$$\mathbb{P}(\tilde{Z}_{i,j}^k > \tilde{Z}_{i,j}^l \text{ for } l \neq k) = \frac{\exp(\alpha_{i,j}^k)}{\sum_{l=1}^A \exp(\alpha_{i,j}^l)}.$$

Thus  $Z_{i,j}^k$  can be obtained by taking arg max operation over  $\tilde{Z}_{i,j}^k$ , i.e.

$$Z_{i,j}^k = \begin{cases} 1 & \text{if } k = \arg \max_{l=1,\dots,A} \{\tilde{Z}_{i,j}^l\} \\ 0 & \text{otherwise} \end{cases} \quad (6)$$


$$\tilde{Z}_{i,j} = [0.1, 0.3, 0.1, 0.4, 0.05, 0.05] \longrightarrow Z_{i,j} = [0, 0, 0, 1, 0, 0] \quad (7)$$

## Proposition 2

$\tilde{Z}_{i,j}^k$  converge to  $Z_{i,j}^k$  in distribution when  $\lambda \rightarrow 0$ , i.e.

$$\mathbb{P}(\lim_{\lambda \rightarrow 0} \tilde{Z}_{i,j}^k = 1) = \frac{\exp(\alpha_{i,j}^k)}{\sum_{l=1}^A \exp(\alpha_{i,j}^l)}.$$

# Gradient estimators: Intuition of Gumbel-Softmax



**Figure:** A visualization of distributions of Gumbel Softmax with different temperatures. Temperature  $\lambda = 0$  correspond to original discrete distribution without relaxation.  $\tilde{Z}_{i,j}^k$  becomes sharper and converge to discrete distribution as  $\lambda \rightarrow 0$ .

## Analysis

- The gradient of the new objective  $\mathbb{E}_U[L_\theta(g(\alpha, U))]$  can be estimated by calculating back-propagation with automatic differentiation libraries
- Biased estimator due to changes of objective
- Low variance thanks to reparameterization

NAS with unbiased and low variance gradient estimators

# Gradient estimators: SF estimators with Control Variates

## Control Variates

In SF estimator, the gradient is estimated using

$$\nabla_{\alpha} \mathbb{E}_{Z \sim p_{\alpha}(Z)}[L_{\theta}(Z)] = \mathbb{E}_{Z \sim p_{\alpha}(Z)}[L_{\theta}(Z) \nabla_{\alpha} \log p_{\alpha}(Z)]. \quad (8)$$

Alternatively, the gradient can also be estimated by

$$\begin{aligned} \nabla_{\alpha} \mathbb{E}_{Z \sim p_{\alpha}(Z)}[L_{\theta}(Z)] = \\ \mathbb{E}_{Z \sim p_{\alpha}(Z)}[(L_{\theta}(Z) - c(Z)) \nabla_{\alpha} \log p_{\alpha}(Z)] + \mathbb{E}_{Z \sim p_{\alpha}(Z)}[c(Z) \nabla_{\alpha} \log p_{\alpha}(Z)], \end{aligned} \quad (9)$$

for some function  $c(Z)$ .

## Analysis

- Lower variance than ordinary SF estimator if  $c(Z)$  is positively correlated with  $L_{\theta}(Z) \nabla_{\alpha} \log p_{\alpha}(Z)$ .
- Require knowledge of  $\mathbb{E}_{Z \sim p_{\alpha}(Z)}[c(Z) \nabla_{\alpha} \log p_{\alpha}(Z)]$

# Gradient estimators: SF estimators with Control Variates

## Constant Control Variates

If we take  $c(Z) = c$  which is constant, then the estimation can be written as

$$\nabla_{\alpha} \mathbb{E}_{Z \sim p_{\alpha}(Z)} [L_{\theta}(Z)] = \mathbb{E}_{Z \sim p_{\alpha}(Z)} [(L_{\theta}(Z) - c) \nabla_{\alpha} \log p_{\alpha}(Z)]. \quad (10)$$

A common used technique is to take  $c$  as the moving average of  $L_{\theta}(Z)$  in each iteration.

We have implemented SF estimator with constant control variates and call it SF in our later experiment.

## Analysis

- Significantly reduce the variance in practice
- Little additional computational cost compare to ordinary version of SF estimator (additional computation is to calculate the moving average)
- Unbiased estimator



# Gradient estimators: RELAX estimators

(Grathwohl et al., 2018) proposed an estimator which they call RELAX, to combine the SF estimator, reparameterization trick and control variates.

## RELAX estimators

RELAX estimators are based on

$$\begin{aligned} \nabla_{\alpha} \mathbb{E}_{Z \sim p_{\alpha}(Z)} [L_{\theta}(Z)] = \\ \mathbb{E}_{U, \tilde{Z}_{cond}} [(L_{\theta}(Z) - c_{\phi}(\tilde{Z}_{cond})) \nabla_{\alpha} \log p_{\alpha}(Z)] + \mathbb{E}_{U, \tilde{Z}_{cond}} [\nabla_{\alpha} (c_{\phi}(\tilde{Z}) - c_{\phi}(\tilde{Z}_{cond}))] \end{aligned} \quad (11)$$

where  $\tilde{Z}$  is Gumbel-Softmax variable,  $\tilde{Z}_{cond} \sim \tilde{Z}|Z$  and  $c_{\phi}$  is a neural network which is trained in each iteration to minimize the variance of the estimator.

# Gradient estimators: Finite-Difference estimators (RAM)


Estimator proposed by (Tokui and sato, 2016), which they call RAM (Reparameterization and Marginalization) estimator

## A simple example

A single binary stochastic variable (i.e. a DAG with only two nodes: input node and output node, with two candidate operations:  $\{O^1, O^2\}$ ). We denote  $e_1 = [1, 0]$  and  $e_2 = [0, 1]$  the  $i$ th entry vector. Then we can calculate the expected loss function as:

$$\mathbb{E}_{Z \sim p_\alpha(Z)}[L_\theta(Z)] = p(e_2)L_\theta(e_2) + (1 - p(e_2))L_\theta(e_1), \quad (12)$$

where  $p(e_1) = \mathbb{P}(Z = e_1)$  and  $p(e_2) = \mathbb{P}(Z = e_2)$ .



# Gradient estimators: Finite-Difference estimators (RAM)

Its gradient w.r.t.  $\alpha$  can be calculated:

$$\nabla_{\alpha} \mathbb{E}_{Z \sim p_{\alpha}(Z)} [L_{\theta}(Z)] = \nabla_{\alpha} p(e_2) (L_{\theta}(e_2) - L_{\theta}(e_1)). \quad (13)$$

## RAM for multi-node and multi-candidate operation

$$\begin{aligned} \nabla_{\alpha} \mathbb{E}_{Z \sim p_{\alpha}(Z)} [L_{\theta}(Z)] &= \sum_{(i,j) \in E} \sum_{Z_{\setminus(i,j)}} p_{\setminus(i,j)}(Z_{\setminus(i,j)}) \times \\ &\sum_{e_a, e_b} (\nabla_{\alpha} \alpha_{i,j}^a) p_{i,j}(e_a) p_{i,j}(e_b) [L_{\theta}(Z_{i,j} = e_a, Z_{\setminus(i,j)}) - L_{\theta}(Z_{i,j} = e_b, Z_{\setminus(i,j)})], \end{aligned} \quad (14)$$

where  $p_{i,j}$  is the marginal probability of  $Z_{i,j}$ ,  $Z_{\setminus(i,j)}$  represent  $Z$  without  $Z_{i,j}$  and  $p_{\setminus(i,j)}$  is the marginal probability of  $Z_{\setminus(i,j)}$ .

# Gradient estimators: Finite-Difference estimators (RAM)

$$\sum_{(i,j) \in E} \sum_{Z_{\setminus(i,j)}} p_{\setminus(i,j)}(Z_{\setminus(i,j)}) \times \\ \sum_{e_a, e_b} (\nabla_{\alpha} \alpha_{i,j}^a) p_{i,j}(e_a) p_{i,j}(e_b) [L_{\theta}(Z_{i,j} = e_a, Z_{\setminus(i,j)}) - L_{\theta}(Z_{i,j} = e_b, Z_{\setminus(i,j)})]. \quad (15)$$

## Sampling strategy

For each edge  $(i,j) \in E$ ,

- sample  $Z_{\setminus(i,j)}$
- calculate

$$\sum_{e_a, e_b} (\nabla_{\alpha} \alpha_{i,j}^a) p_{i,j}(e_a) p_{i,j}(e_b) [L_{\theta}(Z_{i,j} = e_a, Z_{\setminus(i,j)}) - L_{\theta}(Z_{i,j} = e_b, Z_{\setminus(i,j)})] \quad (16)$$

for each  $e_a, e_b$  from candidate operation

# Gradient estimators: Finite-Difference estimators (RAM)

## Analysis

- Unbiased estimator
- Very low variance due to  $L_{\theta}(Z_{i,j} = e_a, Z_{\setminus(i,j)}) - L_{\theta}(Z_{i,j} = e_b, Z_{\setminus(i,j)})$  are evaluated at the same  $Z_{\setminus(i,j)}$
- Theoretically proved to be better (lower variance) than Score Function estimator with constant control variates (Tokui and sato, 2016)
- More computations compare to SF and RELAX

# Gradient estimators conclusion


We have proposed 3 different estimators:

- Score Function estimators with constant control variates (SF)
- RELAX estimators (RELAX)
- RAM estimators (RAM)

## Experiments

# Experiments setting


- For a better comparison, use exactly the same setting as the other research papers: NAS, ENAS, DARTS, SNAS, etc.
- Search for two kinds of cells: Normal cell and Reduction cell
- Implement three kinds of strategies: SF with constant control variates, RELAX and RAM




**Figure:** A conceptual of global structure. where reduction cells are located in  $1/3$  and  $2/3$  of the total depth of the neural network.



# Results



(a) Normal cell found by SF




(b) Reduction cell found by SF


**Figure:** Normal cell and reduction cell (child graph) found by SF on CIFAR-10.

(a) Normal cell. (b) Reduction cell.

# Results




(a) Normal cell found by RELAX after restrain




(b) Reduction cell found by RELAX after restrain

**Figure:** Normal cell and reduction cell (child graph) found by RELAX on CIFAR-10. (a) Normal cell. (b) Reduction cell.

# Results




(a) Normal cell found by RAM




(b) Reduction cell found by RAM

Figure: Normal cell and reduction cell (child graph) found by RAM on CIFAR-10.  
(a) Normal cell. (b) Reduction cell.

# Results



(a) Validation accuracy during architecture search with RAM, RELAX and Score Function estimator.



(b) Entropy of architecture distribution (i.e.  $Z \sim p_{\alpha}(Z)$ ) during architecture search with RAM, RELAX and Score Function estimator.

Figure: Validation accuracy and entropy of architecture distribution during architecture search

# Comparison with other methods

Architecture	Test Error(%)	Params(M)	Search Cost (GPU days)	Search Method
NASNet-A [Zoph and Le, 2016]	2.65	3.3	1800	RL
AmoebaNet-A [Real et al., 2019]	3.34	3.2	3150	evolution
AmoebaNet-B [Real et al., 2019]	2.55	2.8	3150	evolution
Hierarchical Evo [Liu et al., 2017]	3.75	15.7	300	evolution
PNAS [Liu et al., 2018]	3.41	3.2	225	SMOB
ENAS [Pham et al., 2018]	2.89	4.6	0.5	RL
DARTS [Liu et al., 2019]	2.76	3.3	1	gradient-based
SNAS [Xie et al., 2019]	2.85	2.8	1.5	gradient-based
RAM(ours)	2.62	3.6	1.25	gradient-based
SF(ours)	2.64	3.4	0.4	gradient-based
RELAX(ours)	2.70	3.6	0.6	gradient-based

**Figure:** Classification errors of different estimators with other state-of-the-art image classifiers on CIFAR-10. All of our experiments are done using a single V100 GPU.

## Conclusion

# Conclusion

- Develop different gradient-based NAS strategies by introducing unbiased and low variance gradient estimators
- Low computational costs like other gradient-based NAS (around 1 GPU days)
- Results outperform other framework on the same search space

# References I

- Grathwohl, Will et al. (2018). “Backpropagation through the Void: Optimizing control variates for black-box gradient estimation”. In: *arXiv:1711.00123v3 [cs.LG]*.
- J. Williams, Ronald (1992). “Simple statistical gradient-following algorithms for connectionist reinforcement learning”. In: *Reinforcement Learning*, pp. 5-32.
- Liu, Chenxi et al. (2018). “Progressive Neural Architecture Search”. In: *arXiv:1712.00559v3 [cs.CV]*.
- Liu, Hanxiao, Karen Simonyan, Oriol Vinyals, et al. (2017). “Hierarchical Representations for Efficient Architecture Search”. In: *arXiv:1711.00436v2 [cs.LG]*.
- Liu, Hanxiao, Karen Simonyan, and Yiming Yang (2019). “DARTS: Differentiable Architecture Search”. In: *arXiv:1806.09055v2 [cs.LG]*.
- Pham, Hieu et al. (2018). “Efficient Neural Architecture Search via Parameter Sharing”. In: *arXiv:1802.03268v2 [cs.LG]*.



## References II

- Real, Esteban et al. (2019). “Regularized Evolution for Image Classifier Architecture Search”. In: *arXiv:1802.01548v7 [cs.NE]*.
- Tokui, Seiya and Issei sato (2016). “Categorical Reparameterization with Gumbel-Softmax”. In: *arXiv:1611.01239v1 [stat.ML]*.
- Xie, Sirui et al. (2019). “SNAS: Stochastic Neural Architecture Search”. In: *arXiv:1812.09926v2 [cs.LG]*.
- Zoph, Barret and Quoc V. Le (2016). “Neural Architecture Search with Reinforcement Learning”. In: *arXiv:1611.01578v2 [cs.LG]*.

*Thanks!*